Session 01 — GNS3 Lab Setup (Ubuntu)

Duration: ~2 hours **OS**: Ubuntu 25.10

Credentials summary (keep handy):

Machine	Username	Password
Alice / Bob	`ubuntu`	`ubuntu`
Darth	`pi`	`raspberry`
Cisco R1 / SW1	—	no login required

Prerequisites: Download everything first

> Do not start the installation until all downloads are complete. These files are large and the lab cannot proceed without them.

Resource	Link	Size
Cisco C7200 (router image)	Download	~45 MB
Cisco C3745 (switch image)	Download	~38 MB
Ubuntu 25.10 Cloud Image	Download	~784 MB
Ubuntu Cloud Init Data	Download	~128 KB
Raspbian Desktop i386 (Darth)	Download	~3.4 GB

Step 1 — Install GNS3

```
sudo add-apt-repository ppa:gns3/ppa
sudo apt update
sudo apt install qemu-system ubridge gns3-gui gns3-server kitty tigervnc-viewer
```

When prompted during installation:

- *"Should non-superusers be able to capture packets?"* → **Yes**
- *"Should non-superusers be able to use GNS3?"* → **Yes**

Add your user to the required groups, then **log out and log back in**:

```
sudo usermod -aG ubridge,libvirt,kvm,wireshark $(whoami)
```

Verify the installation:

```
gns3server --version
```

Step 2 — Configure NAT (internet access for the lab)

GNS3 uses `libvirt`'s virtual bridge (`virbr0`, `192.168.122.1`) to give lab VMs internet access via NAT.

Do this once — make services start automatically at boot:

```
sudo systemctl enable --now libvirtd.socket
sudo systemctl enable --now virtlogd.socket

sudo virsh net-autostart default
```

Find your internet-facing network interface:

```
ip route | grep '^default' | awk '{print $5}'
```

Example output: `enp0s3` — use your actual interface name in the steps below.

Set up NAT rules and persist them across reboots (replace `enp0s3` with your interface):

```
sudo iptables -t nat -A POSTROUTING -o enp0s3 -j MASQUERADE
sudo iptables -A FORWARD -m conntrack --ctstate RELATED,ESTABLISHED -j ACCEPT
sudo iptables -A FORWARD -i virbr0 -o enp0s3 -j ACCEPT
sudo apt install -y iptables-persistent
sudo netfilter-persistent save
```

Start the virtual bridge now (first time only — after reboot it starts automatically):

```
sudo systemctl is-active libvirtd.socket || sudo systemctl start libvirtd.socket
sudo virsh net-start default
```

If you see `Network default already active` that is fine — continue.

Enable IP forwarding (also needed after reboot until you add it to `/etc/sysctl.conf`):

```
sudo sysctl -w net.ipv4.ip_forward=1
```

> To avoid running the sysctl command after every reboot:
>
> ```
> echo "net.ipv4.ip_forward=1" | sudo tee -a /etc/sysctl.conf
> ```
>
> **Note**: a convenience script that runs all of Step 2 in one command will be published on the course GitHub repository at the end of the session.

Step 3 — First GNS3 launch and server setup

Launch GNS3:

```
gns3
```

On first launch, the Setup Wizard appears:

1. Select "Run the topologies on my local computer":

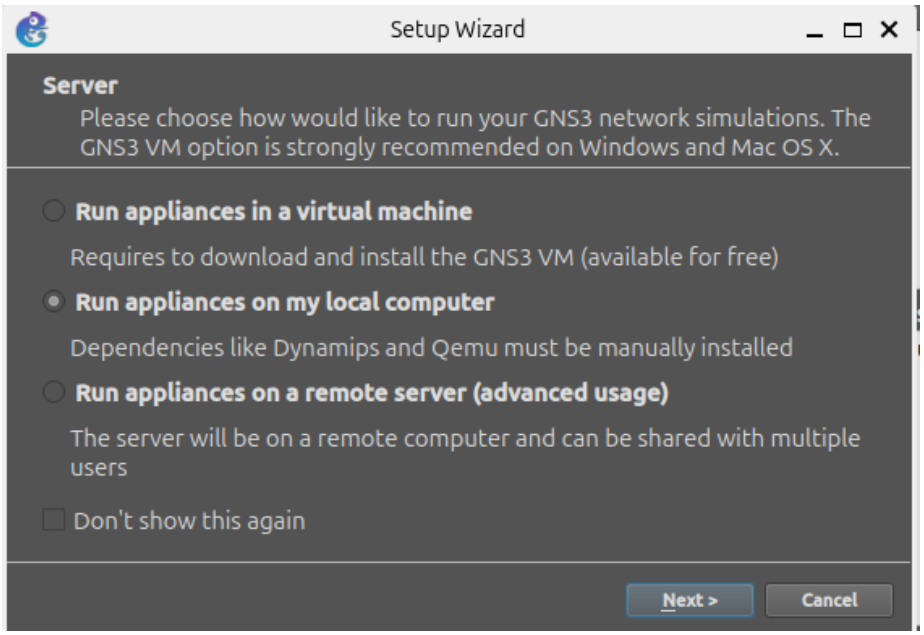

Figure 1: Setup Wizard — server type

2. Accept default local server settings (`localhost:3080`) → Next:

3. Connection successful:

4. Review summary → Finish:

5. GNS3 main window opens:

6. Configure terminal and VNC viewer:

Go to **Edit → Preferences → General → Console applications** and set:

- **Terminal command**: `kitty -e "%c" %i %h %p`
- **VNC viewer**: `vncviewer %h:%p`

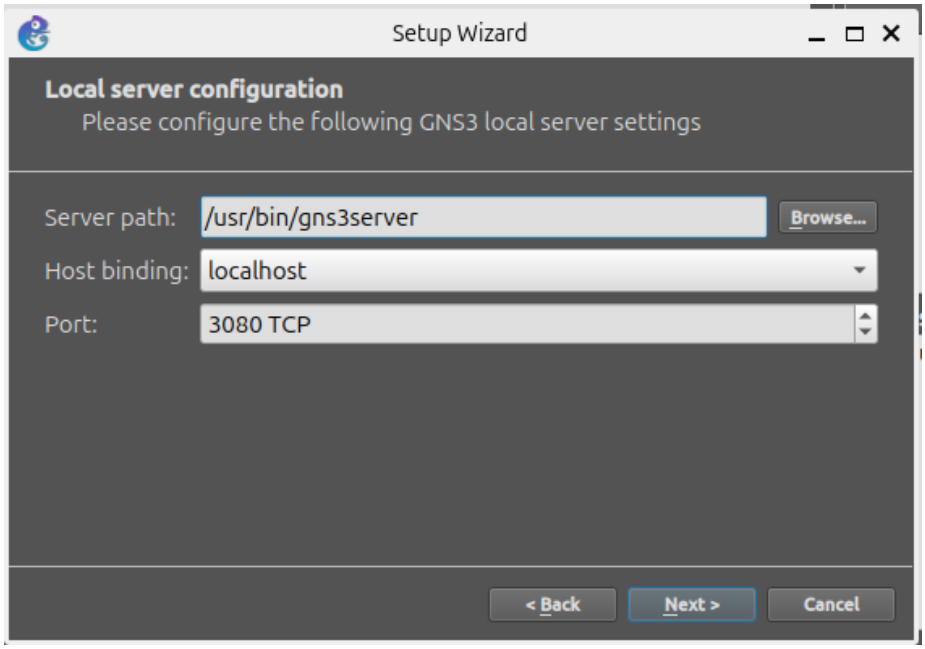

Figure 2: Setup Wizard — local server config

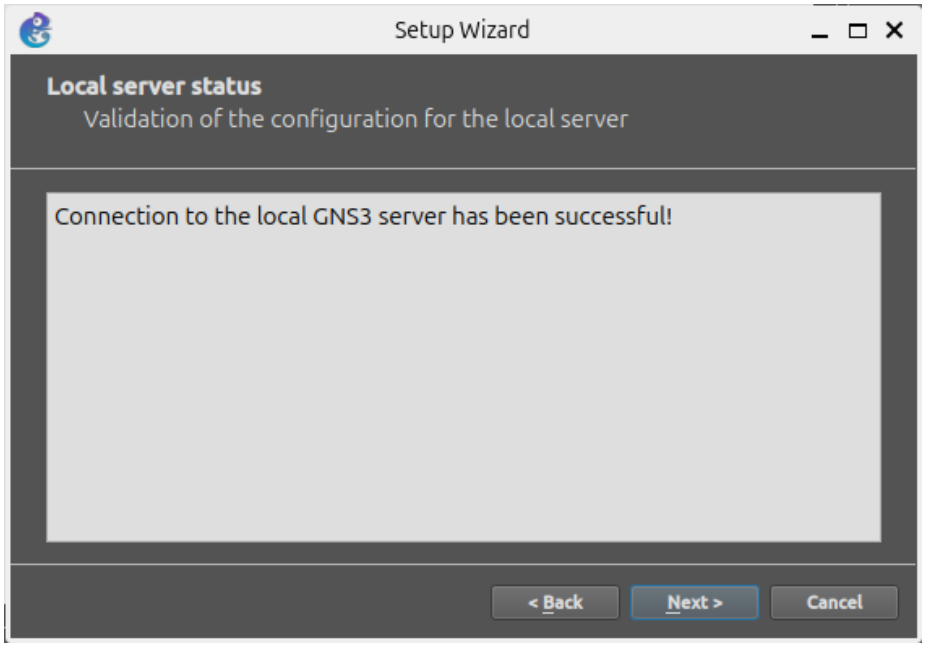

Figure 3: Setup Wizard — connection OK

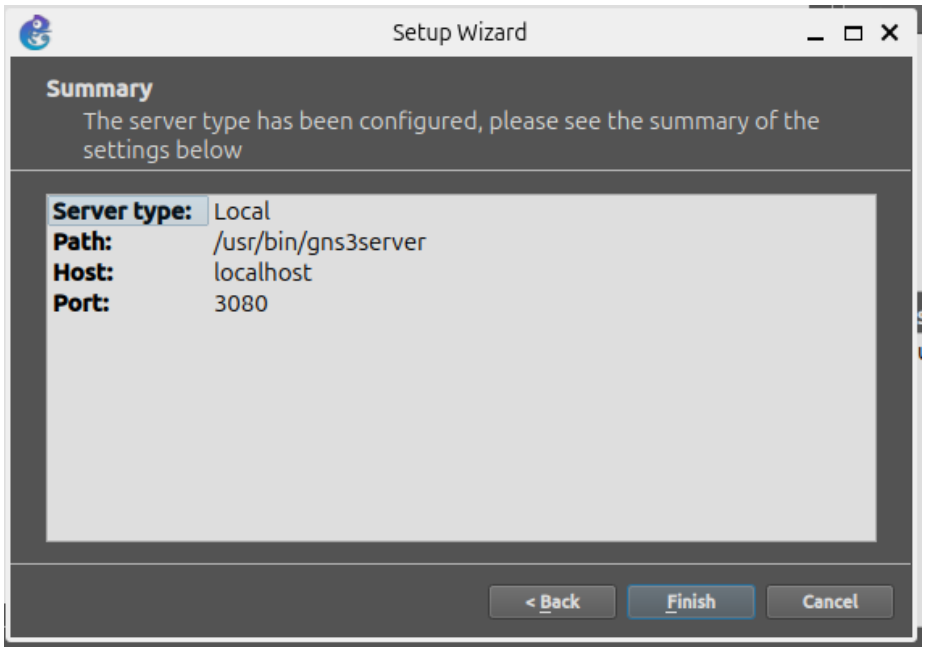

Figure 4: Setup Wizard — summary

`vncviewer` is the binary installed by the `tigervnc-viewer` package — the name is the same.

Step 4 — Create device templates

Templates let you reuse devices across all projects. We need four.

4a. Switch: Cisco C3745

1. Go to **File → New Template**
2. Select "Install an appliance from the GNS3 server"
3. Browse and select the `c3745` image file. If you do not see it, click **Update from online registry** and retry.
4. In the **General** tab:
   - **Symbol**: `:/symbols/affinity/square/blue/switch.svg`
   - **Default name format**: `S{0}`
   - **Initial private-config**: browse to `C3745_startup-config.cfg` from the download folder
5. In the **Slots** tab:
   - `slot 0` → `GT96100-FE`
   - `slot 1` → `NM-16ESW` ← this turns it into a switch

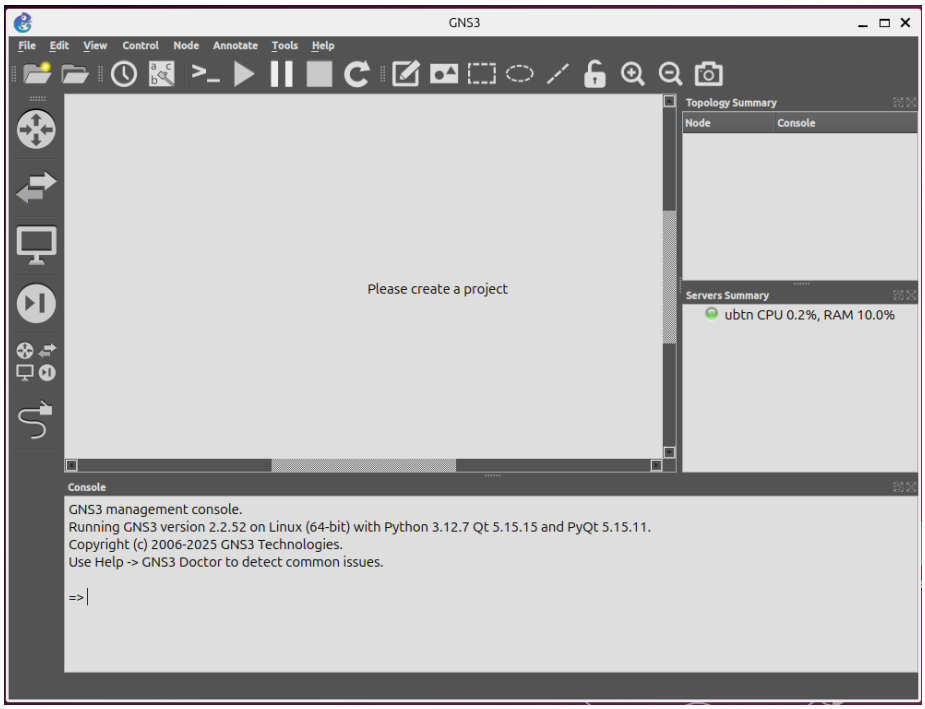

Figure 5: GNS3 main window

6. Click **OK**

   Important: When connecting devices, use ports **FastEthernet1/0 through 1/15** (switch ports). FastEthernet0/0 and 0/1 are routing interfaces — do not use them.

4b. Router: Cisco C7200

1. Go to **File → New Template**
2. Select "Install an appliance from the GNS3 server", browse and select the `c7200` image file. If you do not see it, click **Update from online registry** and retry.
3. In the **General** tab:
   - **Symbol**: `:/symbols/affinity/square/blue/router.svg`
   - **Initial private-config**: browse to `c7200_startup-config.cfg` from the download folder
4. In the **Slots** tab:
   - `slot 0` → `C7200-IO-2FE`
5. Click **OK**

   The router is pre-configured with:

   - `FastEthernet0/0` → NAT outside (`192.168.122.x`)
   - `FastEthernet1/0` → LAN inside (`10.0.0.1`), DHCP server for `10.0.0.0/24`

4c. Ubuntu Cloud Server (Alice and Bob)

The Ubuntu Cloud image acts as a **read-only template**: each VM instance (Alice, Bob) gets its own separate disk that uses the template as a backing file and grows only as data is written, up to the template's maximum size. To set the maximum size for all future instances:

```
qemu-img resize ~/GNS3/images/QEMU/ubuntu-25.10-server-cloudimg-amd64.img +10G
```

> This does not affect instances already created. To resize an existing instance, use `qemu-img resize` on its individual disk file under `~/GNS3/projects/<uuid>/project-files/qemu/<uuid>/hda_disk.qcow2`, then run `sudo growpart /dev/vda 1 && sudo resize2fs /dev/vda1` inside the VM.

Then create the template:

1. Go to **File → New template → Install an appliance from the GNS3 server**
2. Search for `ubuntu cloud` and select **Ubuntu Cloud Guest**
3. Select the base version, then click **Create new version**:
4. Enter the version name `25.10`:

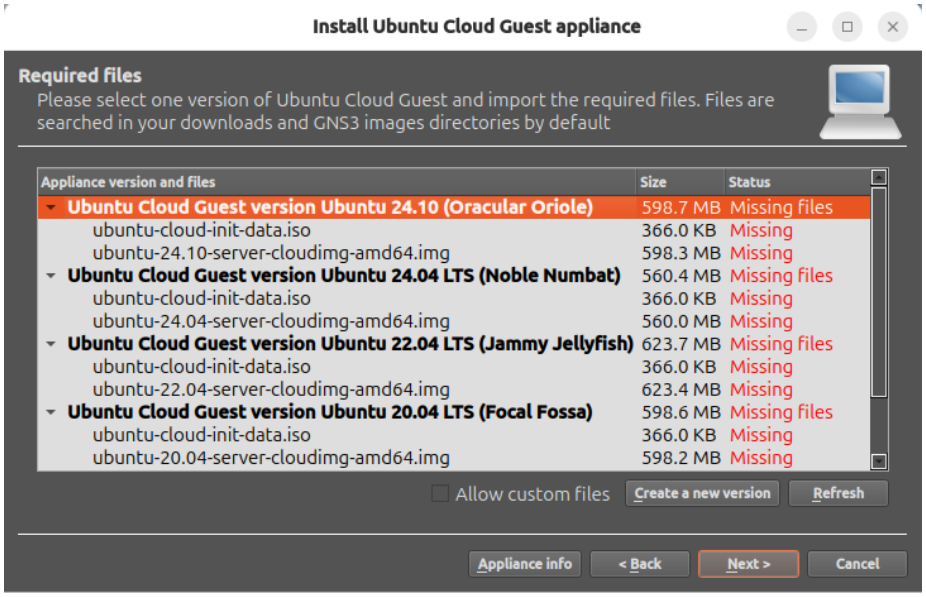

Figure 6: Select base version then click on Create new version

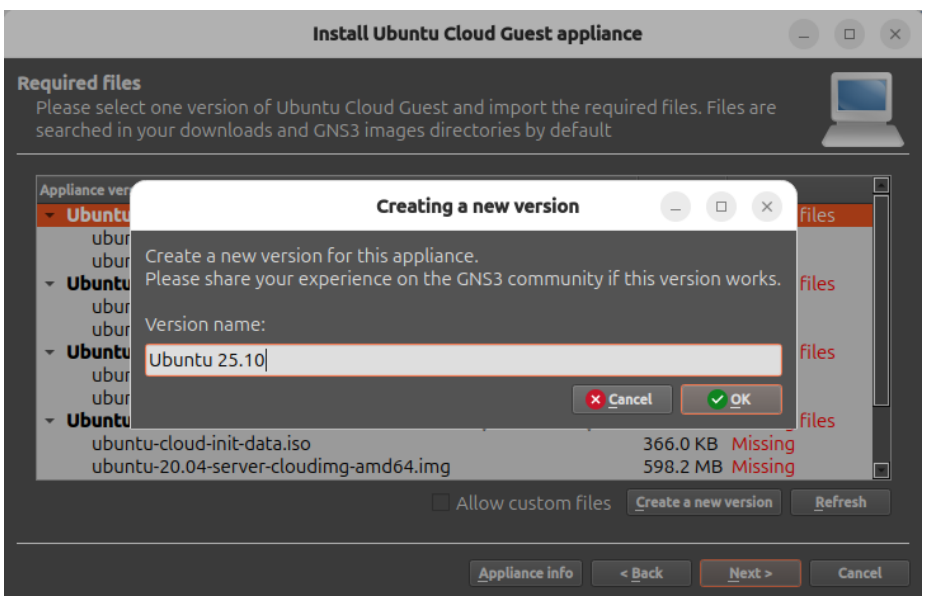

Figure 7: Insert new name

5. Enter the filename of the **cloud-init ISO** (`ubuntu-cloud-init-data.iso`):

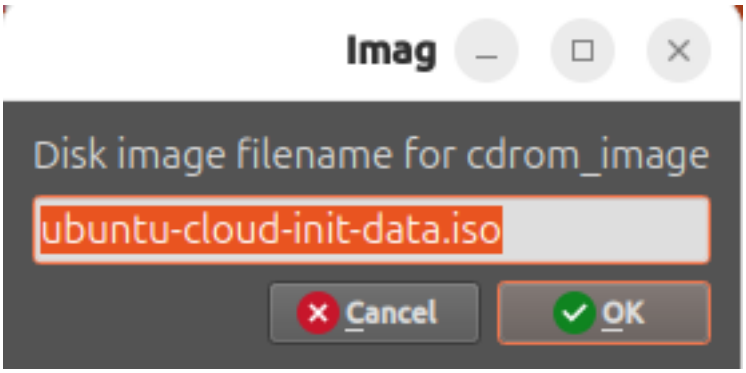

Figure 8: Confirm name of cloud-init ISO image

6. Enter the filename of the **cloud image** (`ubuntu-25.10-server-cloudimg-amd64.img`):

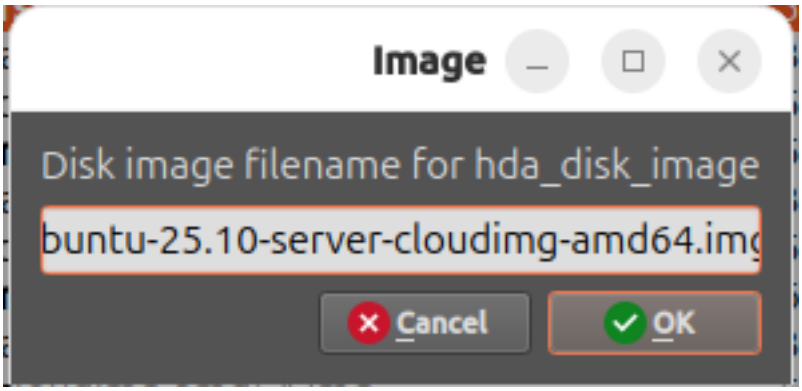

Figure 9: Insert image name

7. The new version is now available — select it and click **Next**:
8. Import the required files:
   - The cloud image (`ubuntu-25.10-server-cloudimg-amd64.img`)
   - The cloud-init ISO (`ubuntu-cloud-init-data.iso`)
9. Click **Finish**

4d. Raspbian (Darth)

Then install the appliance from the GNS3 registry:

1. Go to **File** → **New template** → **Install an appliance from the GNS3 server**

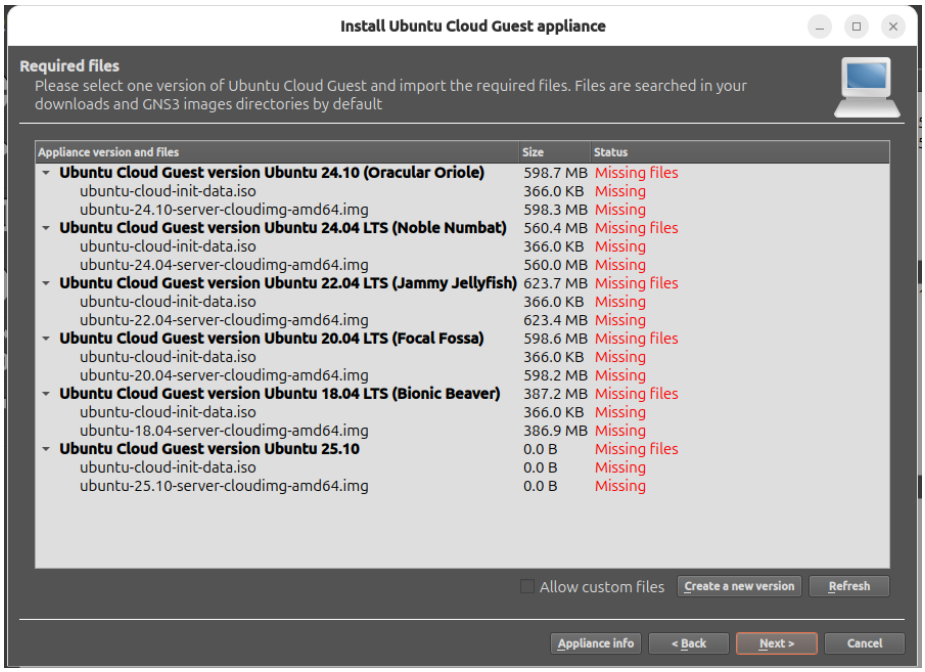

Figure 10: New version available to create template

2. Search for `ras` and select **Raspian** (Qemu / Raspberry Pi Foundation) → **Install**
3. Select **"Install the appliance on the main server"** → **Next**
4. Enable **Allow custom files**, select **Raspberry Pi Desktop version 2022-07-01**, click **Import** next to the ISO and select `2022-07-01-raspios-bullseye-i386.iso` → **Next** → **Finish**

The template is now created but uses a small default disk. You need to replace it with a larger one:

5. Right-click the **Raspian** template in the device panel → **Configure template**
6. Go to the **HDD** tab → next to **HDA Disk image**, click **Create...**
7. In the Qemu image creator, select format **Qcow2** → **Next**:
8. Leave Qcow2 options as default → **Next**:
9. Set file location to `Raspian 2022-07-01-hda.qcow2` and disk size to **30,000 MiB** → **Finish**:
10. The new disk now appears in the **HDA Disk image** field → click **OK** to save the template:

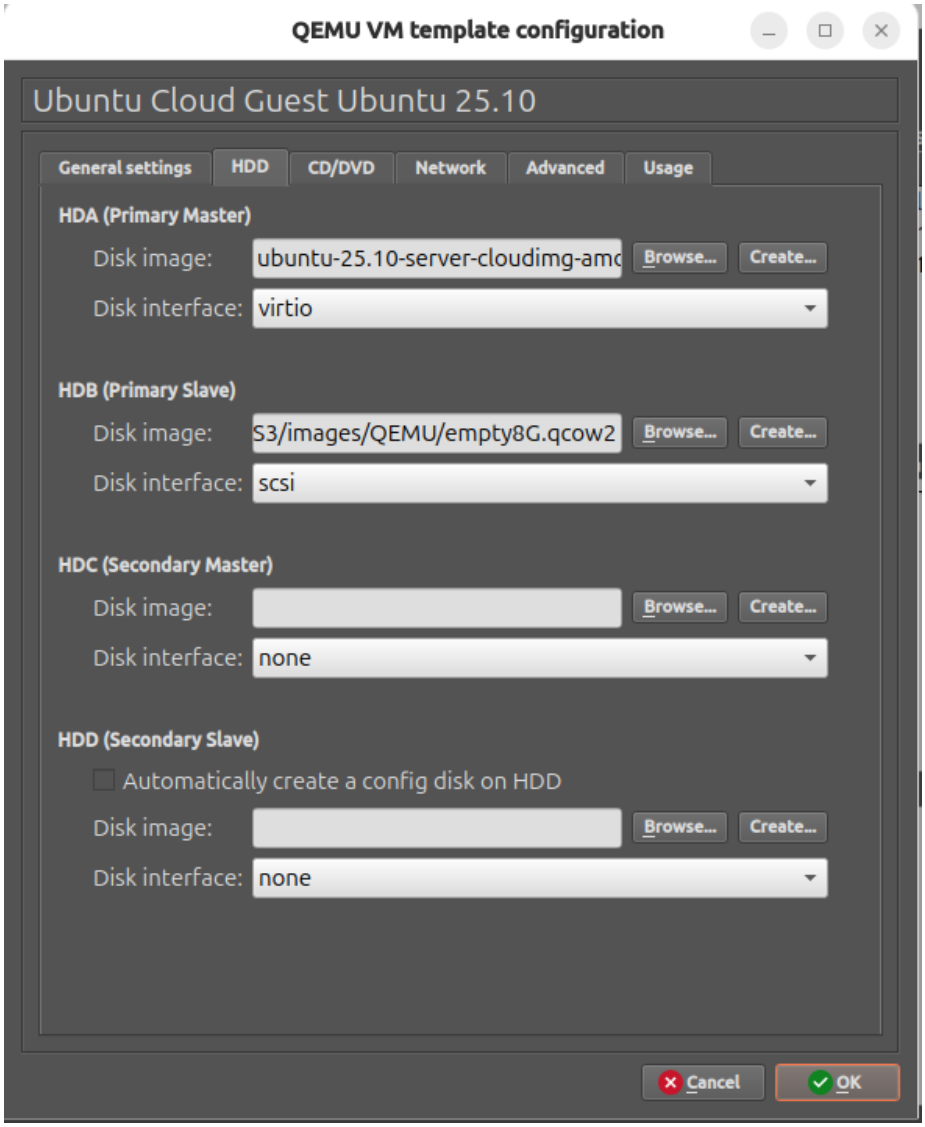

Figure 11: Add disk image

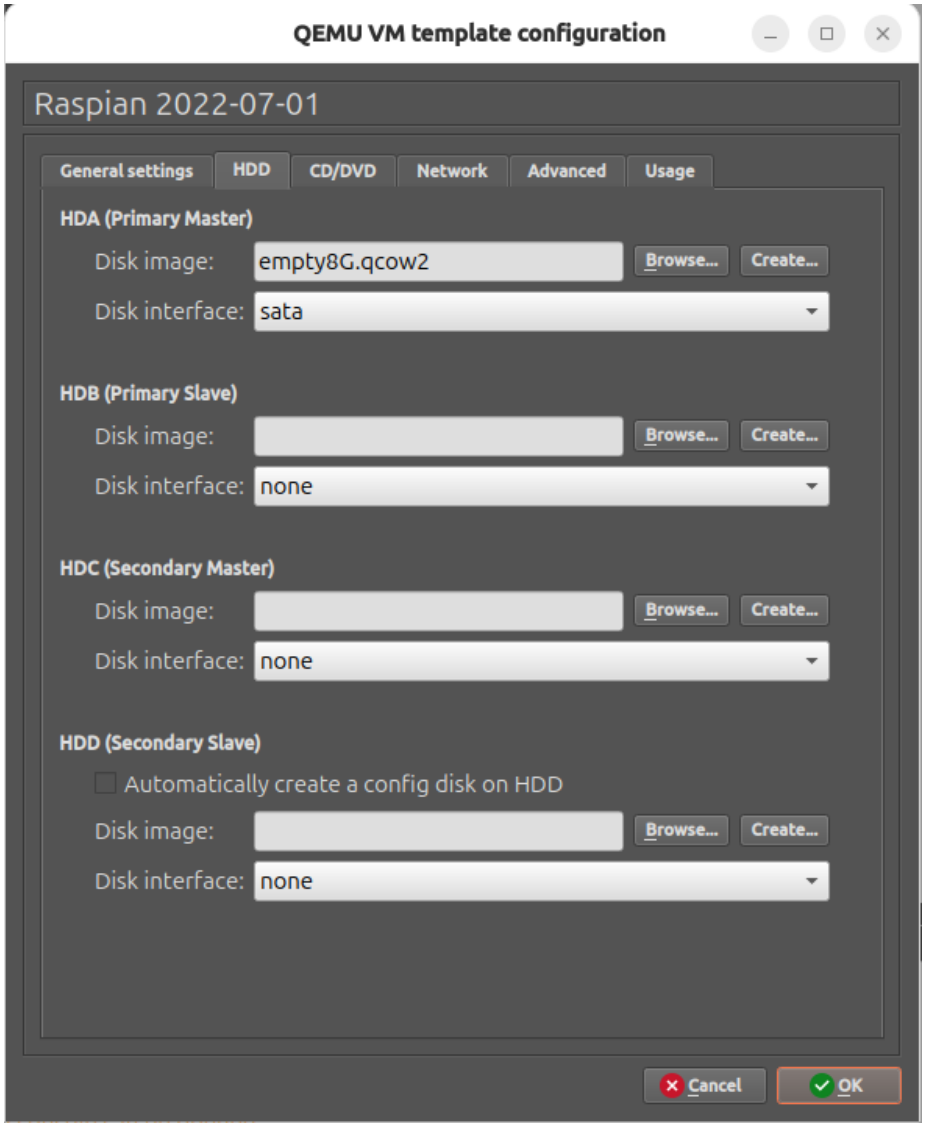

Figure 12: Template in device panel

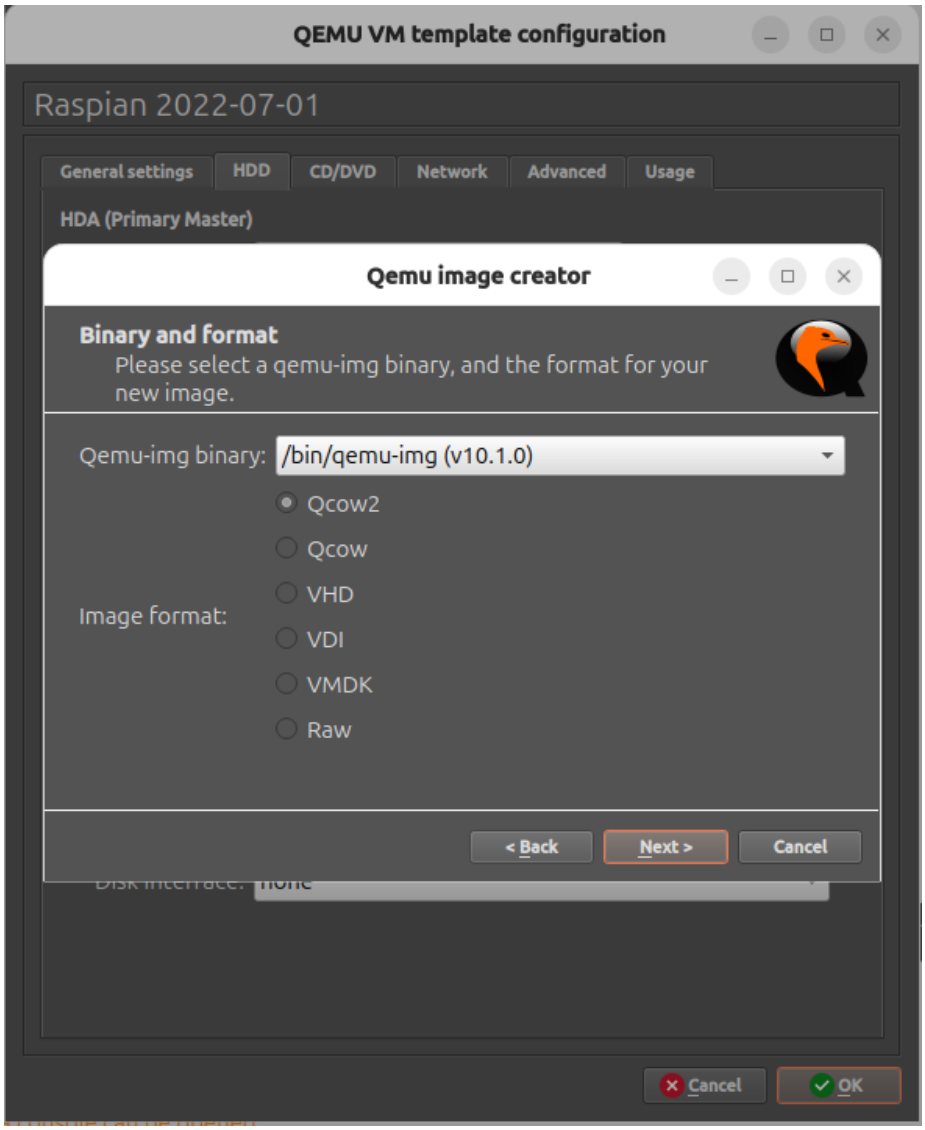

Figure 13: Qemu image creator — format

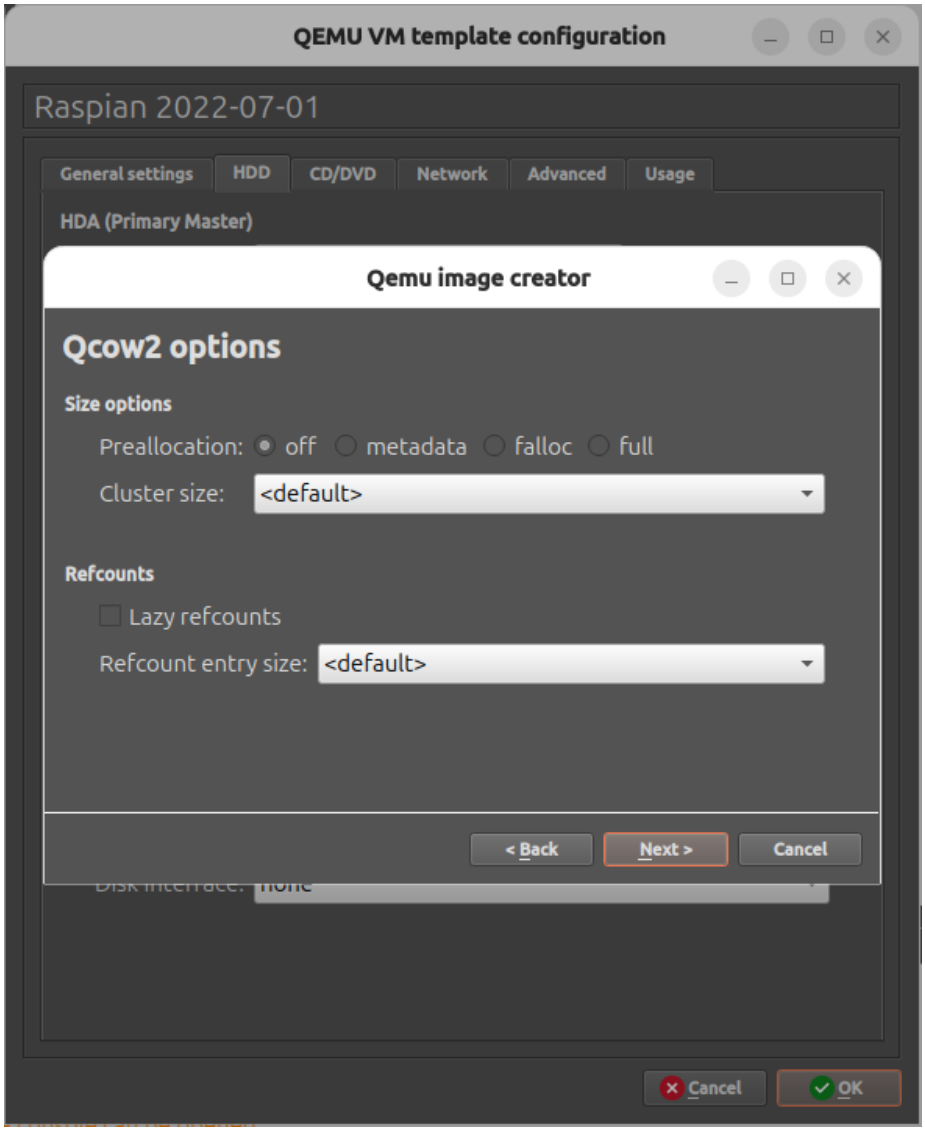

Figure 14: Qemu image creator — options

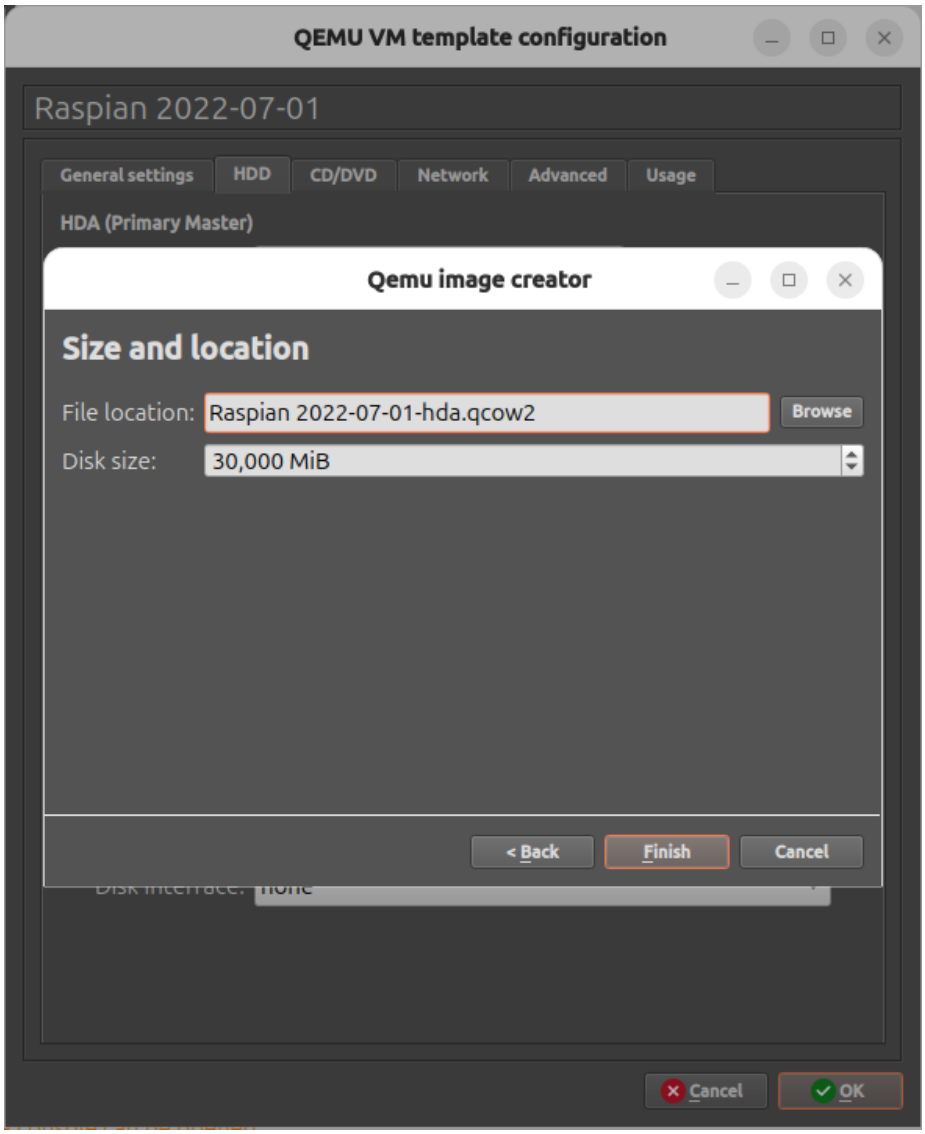

Figure 15: Qemu image creator — size and location

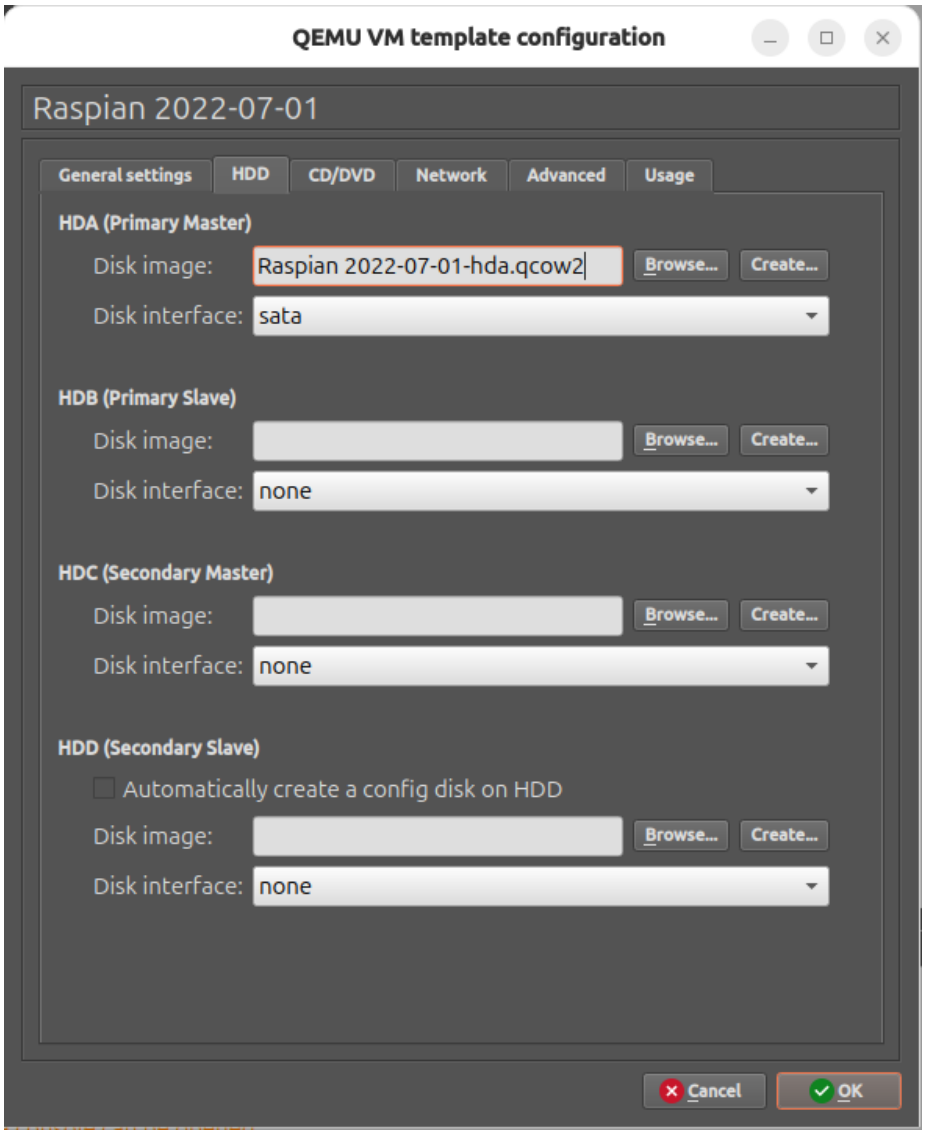

Figure 16: HDD tab with new disk

> Raspbian must be **installed** the first time it boots — it does not work live. See Step 7.

Step 5 — Build the lab topology

Create a new project and add these devices:

IP addresses: | Host | IP | |——|—-| | R1 (gateway) | 10.0.0.1 | | Alice | 10.0.0.10 | | Bob | 10.0.0.20 | | Darth | 10.0.0.30 |

Connect the Cloud node to R1

1. Add a **Cloud** node to the canvas
2. Right-click the Cloud node → **Configure**
3. In the **Ethernet interfaces** list, select `virbr0` → click **Add** → click **OK**

   > If you skip the "Add" step the interface is not actually applied even if it appears selected.

4. Connect Cloud → `R1 FastEthernet0/0`
5. Connect `R1 FastEthernet1/0` → `SW1 FastEthernet1/0`
6. Connect `SW1 FastEthernet1/1` → Alice
7. Connect `SW1 FastEthernet1/2` → Bob
8. Connect `SW1 FastEthernet1/3` → Darth

Step 6 — Configure Alice and Bob (static IP via Netplan)

Login with `ubuntu` / `ubuntu`.

```
sudo nano /etc/netplan/00-lab.yaml
```

For **Alice** (`10.0.0.10`):

```yaml
network:
  version: 2
  renderer: networkd
  ethernets:
    ens3:
      dhcp4: no
      addresses:
        - 10.0.0.10/24
      routes:
        - to: default
```

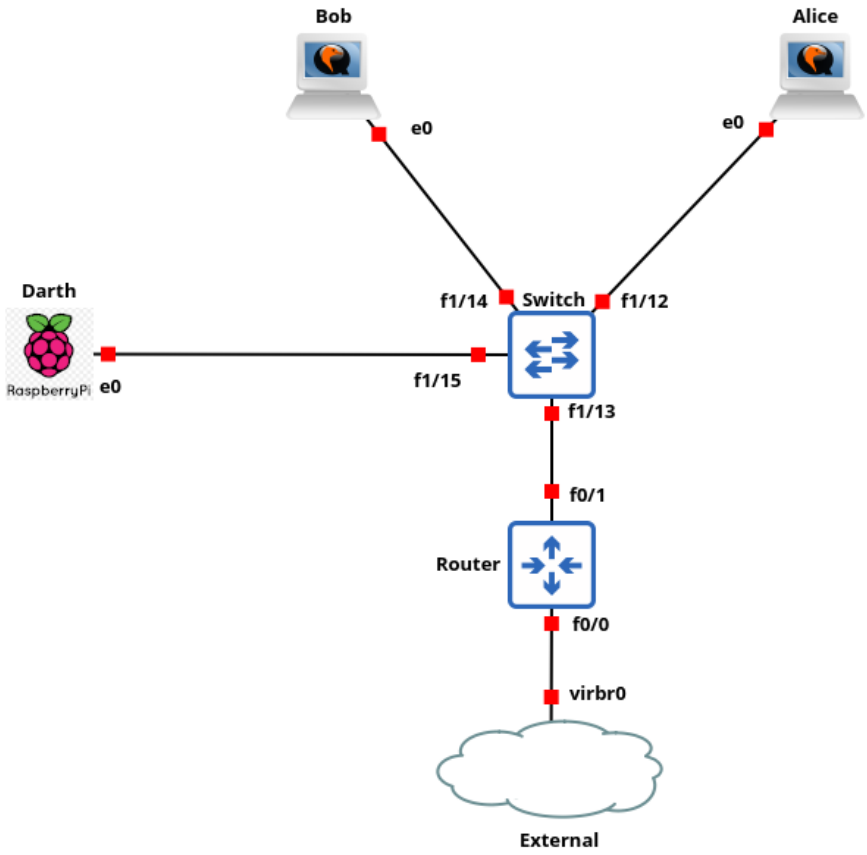

Figure 17: topology

```
        via: 10.0.0.1
    nameservers:
      addresses: [8.8.8.8, 8.8.4.4]
```

For **Bob** (`10.0.0.20`): same file, change address to `10.0.0.20/24`.

Apply on each:

```
sudo chmod 600 /etc/netplan/00-lab.yaml
sudo rm /etc/netplan/50-cloud-init.yaml
sudo netplan apply
ip addr show ens3
resolvectl status   # verify DNS: 8.8.8.8 8.8.4.4
```

then update system with

```
sudo apt update && sudo apt upgrade
```

Step 7 — Install Raspbian OS (first boot only)

Drag the **Raspian** template into the project canvas and start it.

The first boot launches the Debian installer — Raspbian does not work live and must be installed.

> The installation takes **10–15 minutes**. The VM console will appear frozen during disk write — this is normal. Do not stop the VM.

1. At the boot menu, select **Install** (not Graphical install — it may crash)
2. Select your keyboard layout (use arrow keys, then Enter)
3. Partition method: select **Guided – use entire disk**
4. Partitioning scheme: select **All files in one partition**
5. Review the partition summary and select **Finish partitioning and write changes to disk** → confirm **Yes**
6. Wait for the installer to complete. The VM will reboot automatically into the installed system.

   This step only needs to be done once. From the next boot, Darth starts directly into Raspbian.

Step 8 — Configure Darth (static IP)

Login with `pi` / `raspberry`.

Raspbian uses classic Debian networking — not Netplan.

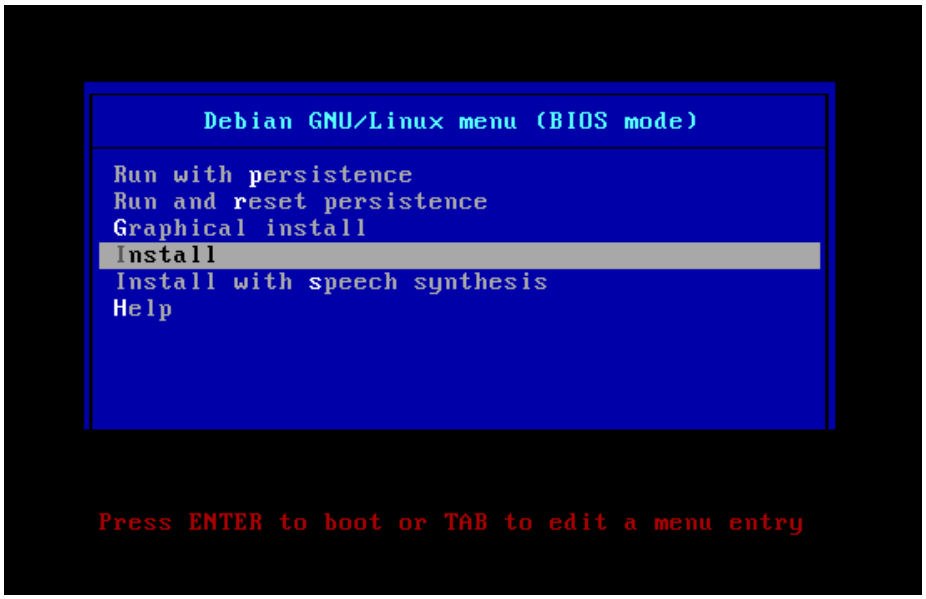

Figure 18: Boot menu

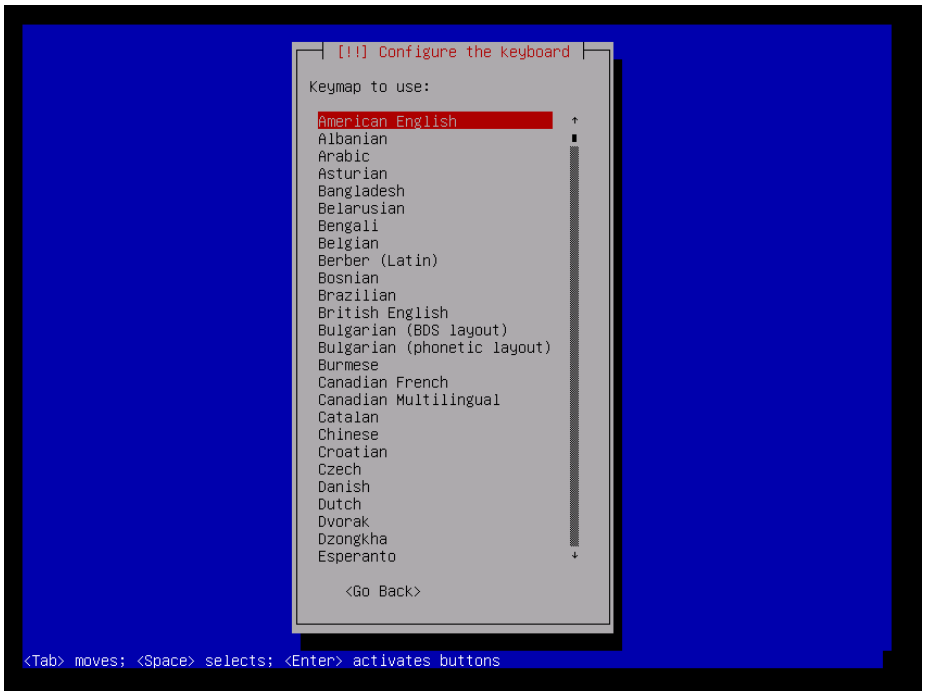

Figure 19: Keyboard layout

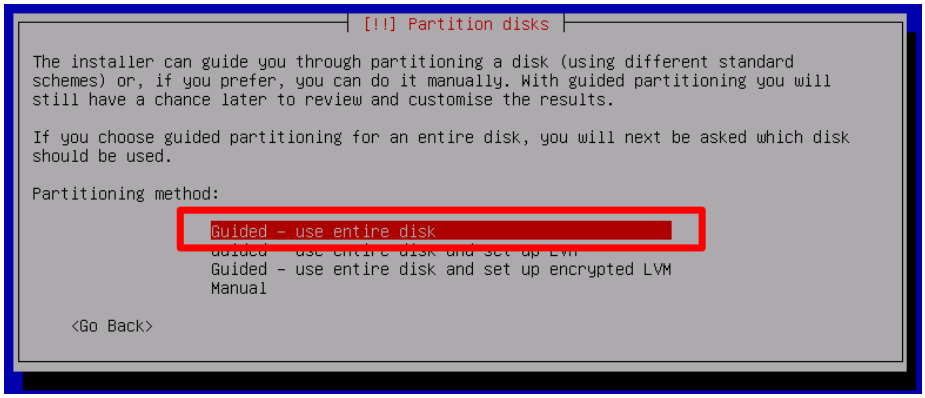

Figure 20: Partition method

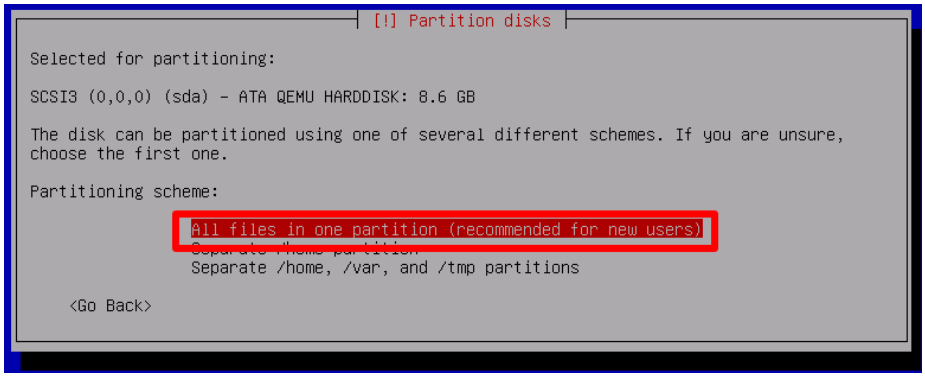

Figure 21: Partition scheme

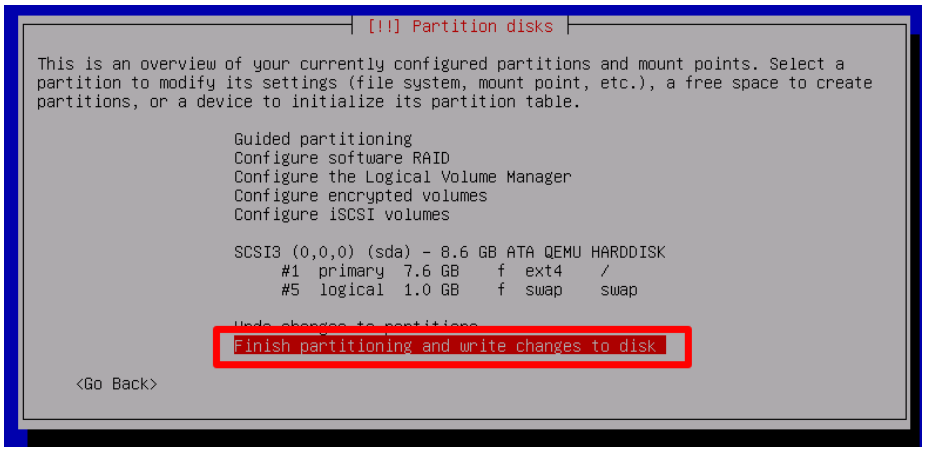

Figure 22: Finish partitioning

```
sudo nano /etc/network/interfaces
```

Add at the end:

```
auto eth0
iface eth0 inet static
    address 10.0.0.30
    netmask 255.255.255.0
    gateway 10.0.0.1
```

Set DNS and lock the file so it is not overwritten by dhclient or resolvconf:

```
echo "nameserver 8.8.8.8" | sudo tee /etc/resolv.conf
echo "nameserver 8.8.4.4" | sudo tee -a /etc/resolv.conf
sudo chattr +i /etc/resolv.conf
```

Apply and verify:

```
sudo systemctl restart networking
ip addr show eth0
cat /etc/resolv.conf
```

Verification checklist

Start all devices (right-click on canvas → Start all nodes).

Alice (`ubuntu` / `ubuntu`):

```
ping -c 3 10.0.0.1       # router
ping -c 3 10.0.0.20      # Bob
ping -c 3 10.0.0.30      # Darth
ping -c 3 8.8.8.8        # internet by IP
resolvectl status         # DNS: 8.8.8.8 8.8.4.4
ping -c 3 google.com     # internet via DNS
```

Bob (`ubuntu` / `ubuntu`):

```
ping -c 3 10.0.0.1       # router
ping -c 3 10.0.0.10      # Alice
ping -c 3 10.0.0.30      # Darth
ping -c 3 8.8.8.8        # internet by IP
resolvectl status         # DNS: 8.8.8.8 8.8.4.4
ping -c 3 google.com     # internet via DNS
```

Darth (`pi` / `raspberry`):

```
ping -c 3 10.0.0.1       # router
ping -c 3 10.0.0.10      # Alice
ping -c 3 10.0.0.20      # Bob
ping -c 3 8.8.8.8        # internet by IP
```

```
cat /etc/resolv.conf      # nameserver 8.8.8.8 / 8.8.4.4
ping -c 3 google.com      # internet via DNS
```

All checks must succeed on every machine.

Troubleshooting

`virbr0` not visible in GNS3 Cloud node

```
sudo systemctl start libvirtd.socket
sudo virsh net-start default
sudo virsh net-autostart default
```

If the `default` network does not exist, create it from scratch using this XML definition:

```
sudo virsh net-define /dev/stdin <<'EOF'
<network>
  <name>default</name>
  <forward mode="nat"/>
  <bridge name="virbr0" stp="on" delay="0"/>
  <ip address="192.168.122.1" netmask="255.255.255.0">
    <dhcp>
      <range start="192.168.122.2" end="192.168.122.254"/>
    </dhcp>
  </ip>
</network>
EOF
sudo virsh net-start default
sudo virsh net-autostart default
```

GNS3 500 error when adding resources (disk space)

Check free space: `df -h ~`

If disk is >90% full, edit `/usr/lib/python3.x/site-packages/gns3server/compute/project_manager.py` and comment out the line `project.emit("log.warning", ...)` inside `_check_available_disk_space`.

`libvirtd` fails: "User record for user 'libvirt-qemu' was not found"

This happens when the `libvirt-qemu` user exists but its home directory is missing or has wrong permissions:

```
sudo mkdir -p /var/lib/libvirt/qemu
sudo chown libvirt-qemu:libvirt-qemu /var/lib/libvirt/qemu
sudo systemctl restart libvirtd.service
```

VM disk full (after the fact)

If you skipped creating the larger disk image, you can expand the per-instance disk after the fact.

Step 1 — Resize the disk file (outside the VM). Close the VM first, then right-click it in GNS3 → **"Show in file manager"** to find the path:

```
qemu-img resize ~/GNS3/projects/<PROJECT_UUID>/project-files/qemu/<VM_UUID>/hda_disk.qcow2 +
```

Step 2 — Expand the partition and filesystem (inside the VM). Resizing the disk file only increases the raw size — the partition and filesystem inside must be expanded too. Boot the VM and run:

For **Ubuntu** (uses `vda`):

```
sudo growpart /dev/vda 1
sudo resize2fs /dev/vda1
df -h /
```

For **Darth** (Raspbian, uses `sda`):

```
sudo growpart /dev/sda 1
sudo resize2fs /dev/sda1
df -h /
```

> To avoid this: always resize the template disk image before creating instances (see Step 4c for Ubuntu, Step 4d for Darth).

GNS3 installation fails: broken `vinagre` dependency

At certain times the `vinagre` package in the GNS3 PPA is broken and blocks installation. Fix it manually:

```
wget http://launchpadlibrarian.net/722482020/vinagre_3.22.0-8ubuntu5_amd64.deb
sudo dpkg -i vinagre_3.22.0-8ubuntu5_amd64.deb
sudo apt-get -f install
```

Then retry the GNS3 installation.

QEMU hardware acceleration compatibility issues

If VMs crash or behave unexpectedly, disable hardware acceleration: go to **Edit → Preferences → QEMU** and uncheck both hardware acceleration options.

`virbr0` uses a different IP than `192.168.122.1`

The router startup config assumes `virbr0` is at `192.168.122.1`. If your libvirt default network uses a different subnet, the router will not get internet access.

Check your virbr0 address:

```
ip addr show virbr0
```

If it differs, either reconfigure the libvirt default network or update the router startup config accordingly.

Terminal columns do not resize automatically

If the terminal inside a VM does not adjust its width when you resize the window:

```
sudo apt install xterm
resize
```

Run `resize` each time you resize the terminal window.